

HA-1

Ishwara Hegde, Jonathan Nieman and Aleksei Samkov

10/3/2020

Preamble

Code chunk loads all the necessary packages and sets the Markdown structure.

Problem 1.3: Lasso vs OLS

Part 1

```
#Calibration
```

```
n<-500  
rho <- 0.5  
p <-11  
B<-500
```

```
#Generating the data
```

```
Sigma <- outer(1:p,1:p,FUN=function(x,y) rho^(abs(x-y)))  
X<-rmvnorm(n,sigma=Sigma,method = "chol")
```

```
beta <- matrix(c(5,rep(1,10)),nrow=p,ncol=1)  
beta_one <- matrix(c(5,rep(1,10)),nrow=11,ncol=1)
```

```
epsilon <-rnorm(n,mean=0,sd=1)
```

```
Y <-X%%beta +epsilon
```

```
#performing OLS B times
```

```
#The "tidy" way to run OLS:
```

```
# I exclude the intercept here cause I include it in definition of beta
```

```
OLS<-lm(Y~0+X)  
OLS<-tidy(OLS)
```

```
beta_hat<-OLS$estimate
```

```
#Estimation error
```

```
error<- t(beta_hat-beta_one)%%solve(cov(X))%%(beta_hat-beta_one)
```

Now I just run a loop to simulate it 500 times. Note, I use the OLS formula below and not the lm model just for variety.

```


est_err<-matrix(,nrow=B,ncol=1)
for (i in 1:B) {
  Sigma <- outer(1:p,1:p,FUN=function(x,y) rho^(abs(x-y)))
  X<-rmvnorm(n,sigma=Sigma,method = "chol")

  epsilon <-rnorm(n,mean=0,sd=1)

  Y <- X%%beta +epsilon
  beta_hat<-solve((t(X)%%X)) %% (t(X)%%Y)
  est_err[i,1]<- t(beta_hat-beta_one)%%cov(X)%%(beta_hat-beta_one)
}

```

####Plot



This resembles a chi-square distribution as expected.

Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

Part 2

```

p<- 101
rho<-0.5
beta<-rbind(beta_one,matrix(0,nrow=p-11,ncol=1))

est_err_ols<-matrix(,nrow=B,ncol=1)

```

```

est_err_lasso<-matrix(,nrow=B,ncol=1)
best_lambdas<-matrix(,nrow=B,ncol=1)
for (i in 1:B) {

  Sigma <- outer(1:p,1:p,FUN=function(x,y) rho^(abs(x-y)))
  X<-matrix(rmvnorm(n,sigma=Sigma,method = "chol"),nrow=500,ncol= p)

  epsilon <-rnorm(n,mean=0,sd=1)

  Y <- X%%beta +epsilon
  beta_hat_ols<-solve((t(X)%*%X)) %% (t(X)%*%Y)
  lasso_model<-cv.glmnet(X,Y,alpha = 1,nfolds = 5,intercept=FALSE)

  #choose the lambda that minimizes MSE
  best_lambdas[i,1] <- lasso_model$lambda.min

  #Now I run LASSO with this lambda


  beta_hat_lasso<- predict(lasso_model, s = best_lambdas[i,1],newx = X,
                           type="coefficients",intercept=FALSE)

  est_err_ols[i,1]<- t(beta_hat_ols-beta)%*%cov(X)%*%(beta_hat_ols-beta)
  est_err_lasso[i,1]<- t(beta_hat_lasso[2:102,]-beta)%*%cov(X)%*%
    (beta_hat_lasso[2:102,]-beta)

}

combined_err<-tibble(ten_dim=(est_err),hund_dim_ols=(est_err_ols)
                     ,hund_dim_lasso = (est_err_lasso))
p<-combined_err %>% ggplot()+
  geom_density(mapping = aes(hund_dim_ols),color="red")+
  geom_density(mapping = aes(hund_dim_lasso),color="green")+
  geom_density(mapping=aes(ten_dim),color="blue")
p+labs(x = "Estimation Error", y = "Density",caption = "red:OLS p=100,
green:lasso, blue:OLS p=10")+theme_minimal()

```



red:OLS p=100,
green:lasso, blue:OLS p=10

From the graph above it seems that the low dimensional OLS performs best. It has lower bias and variance than the other two. However, the fact that LASSO has a lower bias and variance indicates that OLS (p=100) might be overfitting in this case.

Another explanation could be that since there is some correlation between features we can drop a few and improve performance. We can check this logic in part 4.

Part 3

Only the recalibration part changes, rest stays the same as part 2

```
#Recalibration
p=201
rho=0.5

# Creating the sequence
seq=matrix(,nrow = p-11,ncol=1)
for (i in 1:p-11) {seq[i,1]=(1/(2^i))}
beta_three<-rbind(beta_one,seq)

#
est_err_ols_three<-matrix(,nrow=B,ncol=1)
est_err_lasso_three<-matrix(,nrow=B,ncol=1)
best_lambdas_two<-matrix(,nrow=B,ncol=1)
for (i in 1:B) {

  Sigma <- outer(1:p,1:p,FUN=function(x,y) rho^(abs(x-y)))
  X<-matrix(rmvnorm(n,sigma=Sigma,method = "chol"),nrow=500,ncol=p)
```

```

epsilon <-rnorm(n,mean=0,sd=1)

Y <- X%%beta_three +epsilon
beta_hat_ols<-solve((t(X)%*%X)) %*% (t(X)%*%Y)
lasso_model<-cv.glmnet(X,Y,alpha = 1,nfolds = 5,intercept=FALSE)

#choose the lambda that minimizes MSE
best_lambdas_two[i,1] <- lasso_model$lambda.min


#Now I run LASSO with this lambda

beta_hat_lasso<- predict(lasso_model, s = best_lambdas[i,1],newx = X,
                        type="coefficients",intercept=FALSE)

est_err_ols_three[i,1]<-
  t(beta_hat_ols-beta_three)%*%cov(X)%*(beta_hat_ols-beta_three)
est_err_lasso_three[i,1]<-
  t(beta_hat_lasso[2:202,]-beta_three)%*%cov(X)%*
  (beta_hat_lasso[2:202,]-beta_three)
}

combined_err_two<-tibble(hund_dim_ols=est_err_ols_three,hund_dim_lasso
                        =est_err_lasso_three)
p<-combined_err_two %>% ggplot()+geom_density(mapping=aes(hund_dim_ols),
                        color="blue")+
  geom_density(mapping = aes(hund_dim_lasso),color="red")
p+labs(x = "Estimation Error", y = "Density",caption = "Red:Lasso; Blue:OLS" )+
  theme_minimal()

```



Red:Lasso; Blue:OLS

Again, the OLS performs worse than LASSO. I speculate it is for the same reason as before.

Part Four

```
#Recalibration
p=201
rho=0

# Creating the sequence
seq=matrix(,nrow = p-11,ncol=1)
for (i in 1:p-11) {seq[i,1]=(1/(2^i))}
beta_three<-rbind(beta_one,seq)

#
est_err_ols_four<-matrix(,nrow=B,ncol=1)
est_err_lasso_four<-matrix(,nrow=B,ncol=1)
best_lambdas_three<-matrix(,nrow=B,ncol=1)
for (i in 1:B) {

  Sigma <- outer(1:p,1:p,FUN=function(x,y) rho^(abs(x-y)))
  X<-matrix(rmvnorm(n,sigma=Sigma,method = "chol"),nrow=500,ncol=p)

  epsilon <-rnorm(n,mean=0,sd=1)

  Y <- X%*%beta_three +epsilon
}
```

```

beta_hat_ols<-solve((t(X)%*%X)) %*% (t(X)%*%Y)
lasso_model<-cv.glmnet(X,Y,alpha = 1,nfolds = 5,intercept=FALSE)

#choose the lambda that minimizes MSE
best_lambdas_three[i,1] <- lasso_model$lambda.min

#Now I run LASSO with this lambda

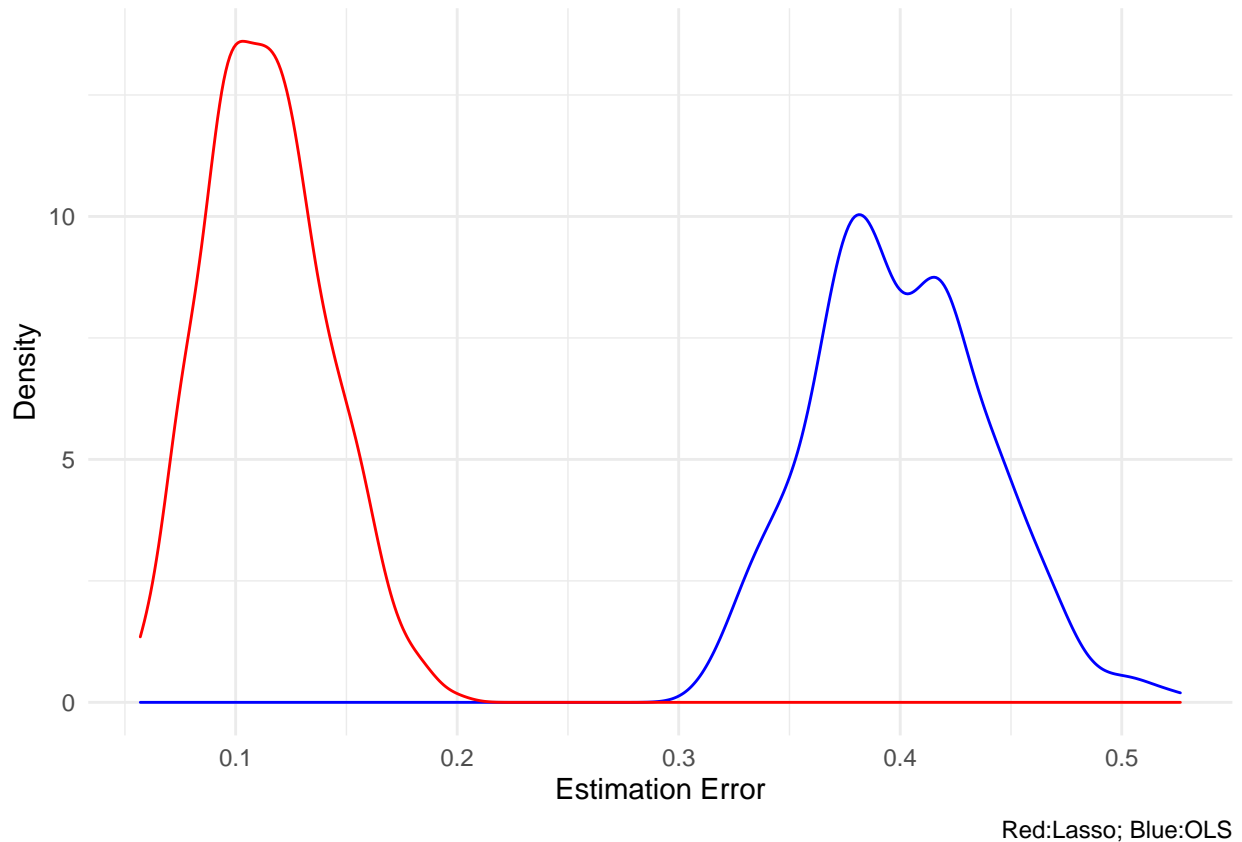
beta_hat_lasso<- predict(lasso_model, s = best_lambdas[i,1],newx = X,
                        type="coefficients",intercept=FALSE)
est_err_ols_four[i,1]<-
  t(beta_hat_ols-beta_three)%*%cov(X)%*(beta_hat_ols-beta_three)

est_err_lasso_four[i,1]<-
  t(beta_hat_lasso[2:202,]-beta_three)%*%cov(X)%*%
  (beta_hat_lasso[2:202,]-beta_three)

}

combined_err_three<-tibble(hund_dim_ols=est_err_ols_four,hund_dim_lasso
  =est_err_lasso_four)
p<-combined_err_three %>% ggplot()+geom_density(mapping=aes(hund_dim_ols),
  color="blue")+
  geom_density(mapping = aes(hund_dim_lasso),color="red")
p+labs(x = "Estimation Error", y = "Density",caption = "Red:Lasso; Blue:OLS" )+
  theme_minimal()

```



What is perhaps a bit surprising is that LASSO performs better even in the 4th case when the features in X are not correlated. It may be then that the penalization is reducing overfitting and not correcting for multicollinearity.